

Introduction to Google Kubernetes

Pascal Liniger, infix development gmbh
Docker Meetup 17.02.15



About me

- Pascal Liniger <pascal.liniger@infix.ch>
- infix development gmbh
- Linux and storage specialist
- Cloud enthusiast
 - Docker
 - Kubernetes
 - etc.

infix

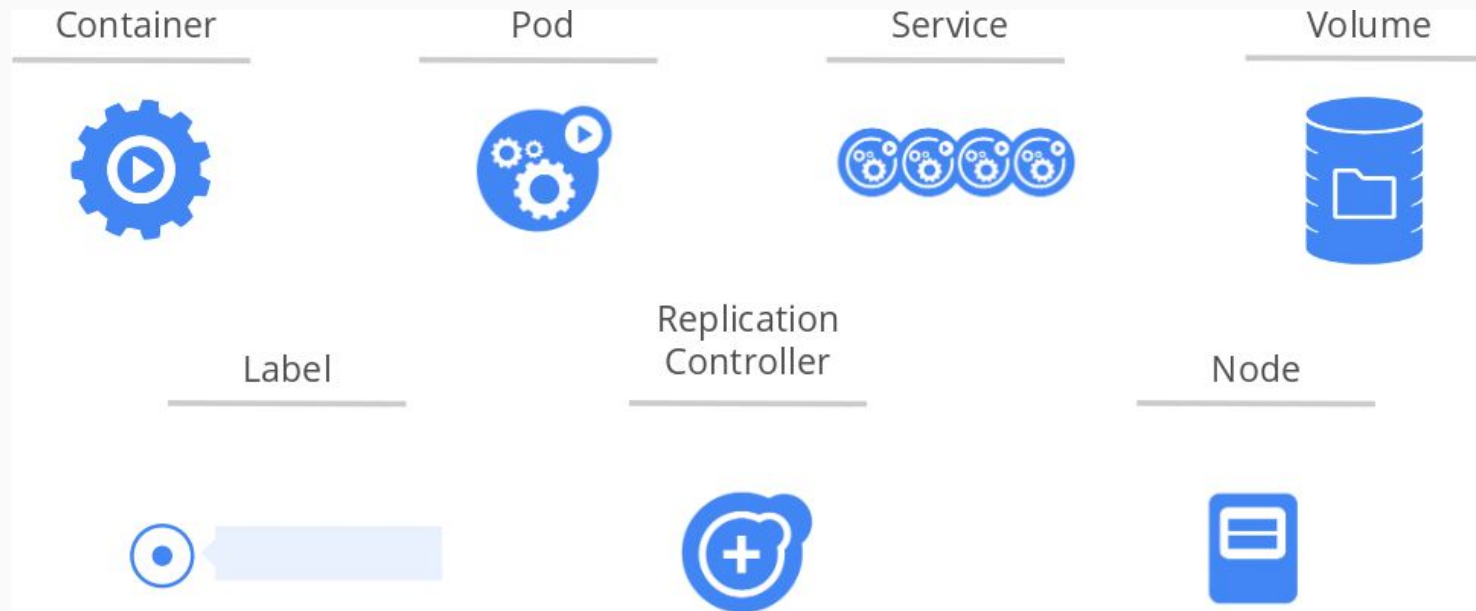
Google Kubernetes

- Greek for “*Helmsman*”
- Orchestrator for **Docker** containers
- Uses Docker to run containers, manage images etc.
- Inspired and based on Google’s internal systems (Borg).
- **Open source**, written in Go




kubernetes

Concepts




Kubernetes Cluster / Nodes



Pods


- Atomic scheduling unit for containers
- **Temporary**: Can die and be replaced
- One pod can contain multiple docker containers
- Containers in pod are tightly coupled
 - Shared IP, namespace, volumes etc.
 - Live and die together

Similar to multiple services running on the same machine



Volumes

- Storage bound to a pod
- Looks like a directory to the containers
- Where and how the data is stored is determined by the **VolumeType**
- Many types available:
 - EmptyDir
 - HostPath
 - nfs
 - glusterfs
 - etc.

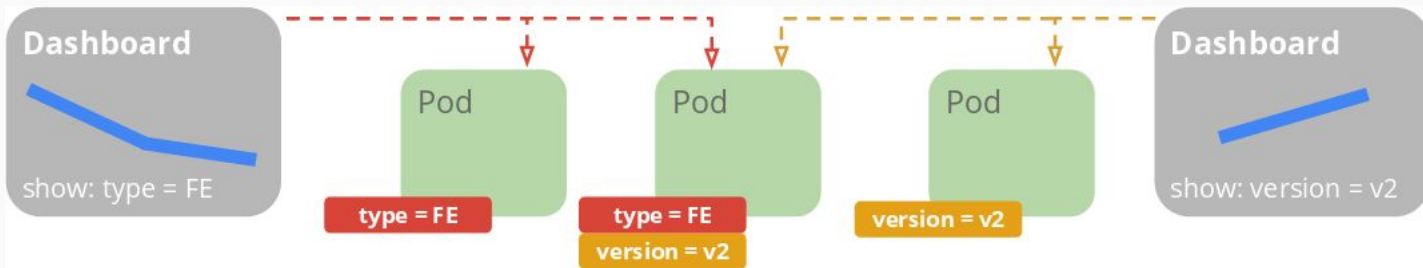


Labels

- Label and group pods
- Used as **selectors** by replication controllers, services and other objects.

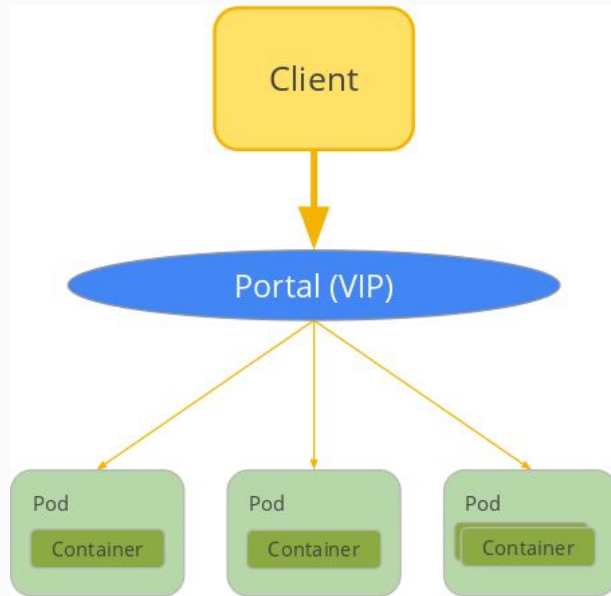
Examples:

- “version”: “v1”, “version”: “v2”
- “release”: “stable”, “release”: “beta”



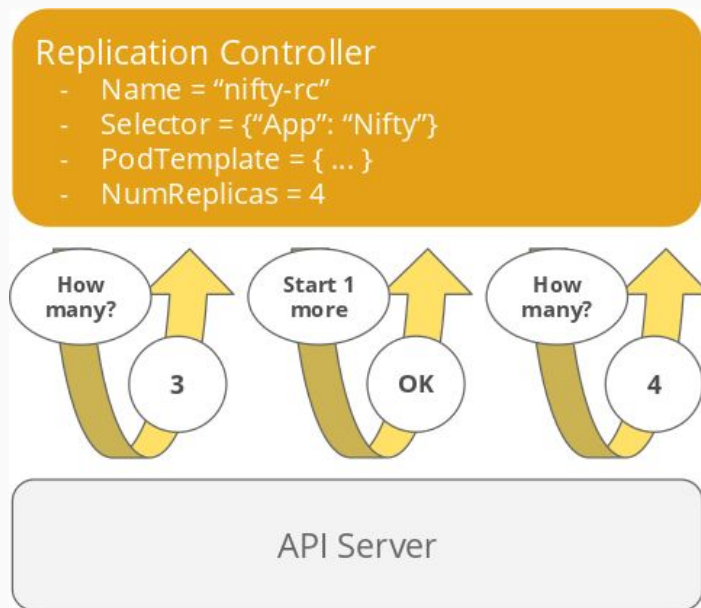
Services

- A group of pods that **act as one**
- Gets a **stable** virtual IP, port and DNS name
- VIP only accessible **inside the cluster**
- Requests get distributed to the pods
- Hides complexity
 - Non-native apps can just access the VIP



Replication Controllers

- **One job:** ensure that N copies of a pod are running
 - If too few, start new ones
 - If too many, kill some
- Selection of pods using **labels**



And more

- Jobs
- Horizontal Pod Autoscaler
- Deployments
- Annotations
- Persistent Volumes and Claims
- Secrets
- etc.

Demo / Cassandra on Kubernetes

- Create a replication controller and a service
- Scale the replication controller
- Distribute a new version of the database (rolling update)
 - While ensuring data availability
 - **Without persistent volumes**
 - Data stays alive by replicating